

# GenreWiz: Music Genre Classification with RNN and CNN

---

[Live Demo](#) | [Github Repo](#)

Due to free Heroku account, please allow a short delay before the webpage starts up

---

## Team

Li Yuxuan 1003607

Keith Ng 1003515

Ng Jia Yi 1003696

---

## Introduction

### Model

### Implementation

Data loader

Training and Validation

Testing

### Results

Accuracy profile

Loss profile

Confusion matrix

Classification Report

Results analysis

### Graphical User Interface

### Conclusion

---

## Introduction

---

The rise in need for automatic music genre classification (AMGC) revolves around the fact that artists today tend to distribute their songs on their websites, making music database management a must. Another recent development is to consume music via streaming, raising the tendency for online radio stations to play songs based on genre preferences. In addition, music browsing and playlist creation these days all utilizes recommendations based on genre.

In this project, we seek to classify music into its respective genre by passing a music spectrogram into a parallel convolutional neural network (CNN) - recurrent neural network (RNN) model.

## Model


---

In our original project proposal, we suggested the usage of a CNN model as we intended to convert audio files into spectrograms - a visual representation of audio across frequency and time dimension. CNN makes perfect sense here since our data inputs into the model are essentially images, and CNN excels at image classification.

On the other hand, RNN excels in understanding sequential data by making the hidden state at time  $t$  dependent on hidden state at time  $t-1$ . The spectrograms have a time component and RNNs can do a much better job of identifying the short term and longer term temporal features in the song.

We referenced the following post from Medium (<https://towardsdatascience.com/using-cnns-and-rnns-for-music-genre-recognition-2435fb2ed6af>) (Priya, 2018).

We implemented the parallel CNN-RNN model as specified in the above post:



Layer (type)	Output Shape	Param #	Connected to
input (InputLayer)	(None, 640, 128, 1)	0	
conv_1 (Conv2D)	(None, 638, 128, 16)	64	input[0][0]
max_pooling2d_1 (MaxPooling2D)	(None, 319, 64, 16)	0	conv_1[0][0]
conv_2 (Conv2D)	(None, 317, 64, 32)	1568	max_pooling2d_1[0][0]
max_pooling2d_2 (MaxPooling2D)	(None, 158, 32, 32)	0	conv_2[0][0]
conv_3 (Conv2D)	(None, 156, 32, 64)	6208	max_pooling2d_2[0][0]
max_pooling2d_3 (MaxPooling2D)	(None, 78, 16, 64)	0	conv_3[0][0]
conv_4 (Conv2D)	(None, 76, 16, 64)	12352	max_pooling2d_3[0][0]
max_pooling2d_4 (MaxPooling2D)	(None, 19, 4, 64)	0	conv_4[0][0]
conv_5 (Conv2D)	(None, 17, 4, 64)	12352	max_pooling2d_4[0][0]
pool_lstm (MaxPooling2D)	(None, 160, 64, 1)	0	input[0][0]
max_pooling2d_5 (MaxPooling2D)	(None, 4, 1, 64)	0	conv_5[0][0]
lambda_1 (Lambda)	(None, 160, 64)	0	pool_lstm[0][0]
flatten_1 (Flatten)	(None, 256)	0	max_pooling2d_5[0][0]
bidirectional_1 (Bidirectional)	(None, 128)	49536	lambda_1[0][0]
concat (Concatenate)	(None, 384)	0	flatten_1[0][0] bidirectional_1[0][0]
preds (Dense)	(None, 8)	3080	concat[0][0]

### Parallel CNN-RNN Model

The convolutional block of the model consists of 2D convolution layer followed by a 2D Max pooling layer. The block is repeated 5 times. The final output is flattened.

The recurrent block starts with 2D max pooling layer of pool size (4,2) to reduce the size of the spectrogram before LSTM operation. This feature reduction was done primarily to speed up processing. The reduced image is sent to a bidirectional GRU with 64 units. .

The outputs from the convolutional and recurrent blocks are then concatenated . Finally we have a dense layer with SoftMax activation.

Here is our PyTorch model :

```

class PCRNN(nn.Module):
    def __init__(self, ks=(3, 1)):
  
```

```

super().__init__()
# CNN
self.conv2d_1 = nn.Conv2d(1, 16, ks)
self.conv2d_2 = nn.Conv2d(16, 32, ks)
self.conv2d_3 = nn.Conv2d(32, 64, ks)
self.conv2d_4 = nn.Conv2d(64, 64, ks)
self.conv2d_5 = nn.Conv2d(64, 64, ks)

self.maxPool2d_1 = nn.MaxPool2d((2, 2))
self.maxPool2d_2 = nn.MaxPool2d((2, 2))
self.maxPool2d_3 = nn.MaxPool2d((2, 2))
self.maxPool2d_4 = nn.MaxPool2d((4, 4))
self.maxPool2d_5 = nn.MaxPool2d((4, 4))

self.flatten_1 = nn.Flatten()

# RNN
self.pool_lstm = nn.MaxPool2d((4, 2))
self.gru = nn.GRU(input_size=64, hidden_size=1, bidirectional=True)

# Overall
self.linear = nn.Linear(576, 10) # 10 genres
self.softmax = nn.Softmax(dim=1)

def forward_cnn(self, x: torch.Tensor):
    x = F.relu(self.conv2d_1(x))
    x = self.maxPool2d_1(x)
    x = F.relu(self.conv2d_2(x))
    x = self.maxPool2d_2(x)
    x = F.relu(self.conv2d_3(x))
    x = self.maxPool2d_3(x)
    x = F.relu(self.conv2d_4(x))
    x = self.maxPool2d_4(x)
    x = F.relu(self.conv2d_5(x))
    x = self.maxPool2d_5(x)
    x = self.flatten_1(x)
    return x

def forward_rnn(self, x: torch.Tensor):
    x = self.pool_lstm(x)
    x = x.view(x.size(0), x.size(2), x.size(3))
    x = self.gru(x)
    x = self.flatten_1(x[0])
    return x

def forward(self, x: torch.Tensor):
    cnn = self.forward_cnn(x)
    rnn = self.forward_rnn(x)
    output = torch.cat([cnn, rnn], dim=1)
    output = self.linear(output)
    output = self.softmax(output)

    return output

```

# Implementation


---

## Data loader

---

We created our customized DataLoader class that is able to read the WAV audio files from `data/` folder and transform them into the Mel-spectrograms in the form of a 2D array for input to the model. Taking into account that audio file length might vary slightly, even though each song sample is about 30 seconds in total, we did some simple cropping of the sample to only take the first 640 units of the spectrogram. This can standardize our input shape to the model as well. The dataloader follows the 8-2 proportion when dividing the dataset into training and validation. Hence, 80% of the samples are used for training and 20% of the samples are used for validation.

An example of a spectrogram looks as such:



We designed the code such that the DataLoader is a simple plug-and-play function that can be used at any stage.

In later development, we also included `test` mode in the same DataLoader. Hence, when switched to `test` mode, this DataLoader will grab the test samples from another directory and output the corresponding 2D spectrogram data as well as the corresponding label.

## Training and Validation

---

We used the Kaggle GTZAN dataset as our training and validation data. The dataset consists of 1000 songs split evenly between 10 music genres. The .wav audio clip (30s snippet) and spectrograms were provided for each song. However, we did not utilize their provided spectrogram as we were unable to replicate it with our own songs.

We split the dataset into 80% training data and 20% validation data. Meaning to say, within each genre, 80 songs were used for training, and 20 songs for validation. The selections were done randomly.

The model was trained using RMSProp optimizer with a learning rate of 0.0005 and the loss function was categorical cross entropy. The model was trained for 120+ epochs and Learning Rate was constantly changed (increased when accuracy plateaued and decreased when accuracy fluctuated too much).

During training, the model that refreshes the highest accuracy will be automatically saved. Hence we make sure that we always obtain the best model after the training.

## Testing

---

For the testing data, we went to YouTube and downloaded 5 songs from each genre of music. In the GTZAN dataset, the songs ranged from the 1970s to the early 2000s. It consisted a mix of recorded songs and live performance. There was also no consistency in the volumes of the songs or the sections of the songs - some were choruses, some were verses, some were instrumentals and some were a mix. Thus, when selecting when preparing the songs for the test data, we ensured that they were as diverse as the ones provided in the GTZAN dataset.

We have an evaluation script which runs the trained model through the 50 test samples we collected and generated a classification report as well as a confusion matrix. The test results are discussed in the next section.


## Results

---

### Accuracy profile

---


Here is the accuracy profile on validation:



### Loss profile


---

Here is the loss profile on train and validation:



## Confusion matrix

We ran our model on 50 test samples that we downloaded from online. Here is the confusion matrix of the model performance. We can see that Classical, Country and Reggae were mostly correctly classified whereas Disco, HipPop, Jazz, Pop and Rock were hard to differentiate.



## Classification Report

Here is the classification report generated on the 50 test samples.

	precision	recall	f1-score	support
blues	0.29	0.40	0.33	5
classical	0.71	1.00	0.83	5
country	0.50	0.80	0.62	5
disco	0.17	0.20	0.18	5
hiphop	0.14	0.20	0.17	5

jazz	0.00	0.00	0.00	5
metal	0.33	0.40	0.36	5
pop	0.00	0.00	0.00	5
reggae	0.50	0.60	0.55	5
rock	0.50	0.20	0.29	5
accuracy			0.38	50
macro avg	0.31	0.38	0.33	50
weighted avg	0.31	0.38	0.33	50
F1: 0.332534				
Recall: 0.380000				
Precision: 0.314286				
Accuracy: 0.380000				

## Results analysis

From the aforementioned results generated, we can see that our model's performance is not fantastic. This can be due to several factors.

Firstly, our training dataset is limited. We only have 1000 songs as input and 800 are used for training and 200 are used for validation. This means that our model is only trained based on the 800 songs, which is not enough.

Secondly, our test samples are found online by ourselves on YouTube, which contain our personal judgement of whether a song belongs to which genre. Hence, our test samples might not be all correctly labelled in the first place. A better strategy if time and resources permit, is to find professional music curator to help to label the test samples to their correct genres. In addition to that, there are no strict boundaries that specify a song genres. Songs could be a mixture of different genres and hence, they might not have one true label.

Therefore, with more training data and more accurately labelled data (multi labelled classification), and perhaps more epochs of training, we foresee an improvement in our model performance.

## Graphical User Interface

We have developed a GUI demo for our Genre Classifier in React, and hosted publicly on Heroku. Due to the free account limitation, after 30 minutes of inactivity, Heroku will shutdown the server. Hence, users may experience a long wait when accessing the website.

Here is the public link to our web demo: <https://genrewiz.herokuapp.com>

In this web demo, we provide the users a simple feature of uploading their WAV audio file and the backend model will predict and output the top 5 possible genres that it identifies.

We have also provided some wav audio samples in the file /sample for you to try out.

Here are the snapshots of our web demo, both on PC and on mobile devices:

# GENREWIZ: YOUR GENRE CLASSIFIER

Upload your music now!



Click or drag your audio file to this area to upload  
Only support .wav (WAV) audio file type.  
Single file upload.

Figure 1: Mobile homepage of GenreWiz



✓ pop1.wav file uploaded successfully.

## GENREWIZ: YOUR GENRE CLASSIFIER

Upload your music now!



Click or drag your audio file to this area to upload  
Only support .wav (WAV) audio file type.  
Single file upload.



pop1.wav



AI is predicting the genres for you...

Figure 2: Upload WAV file successfully and start prediction

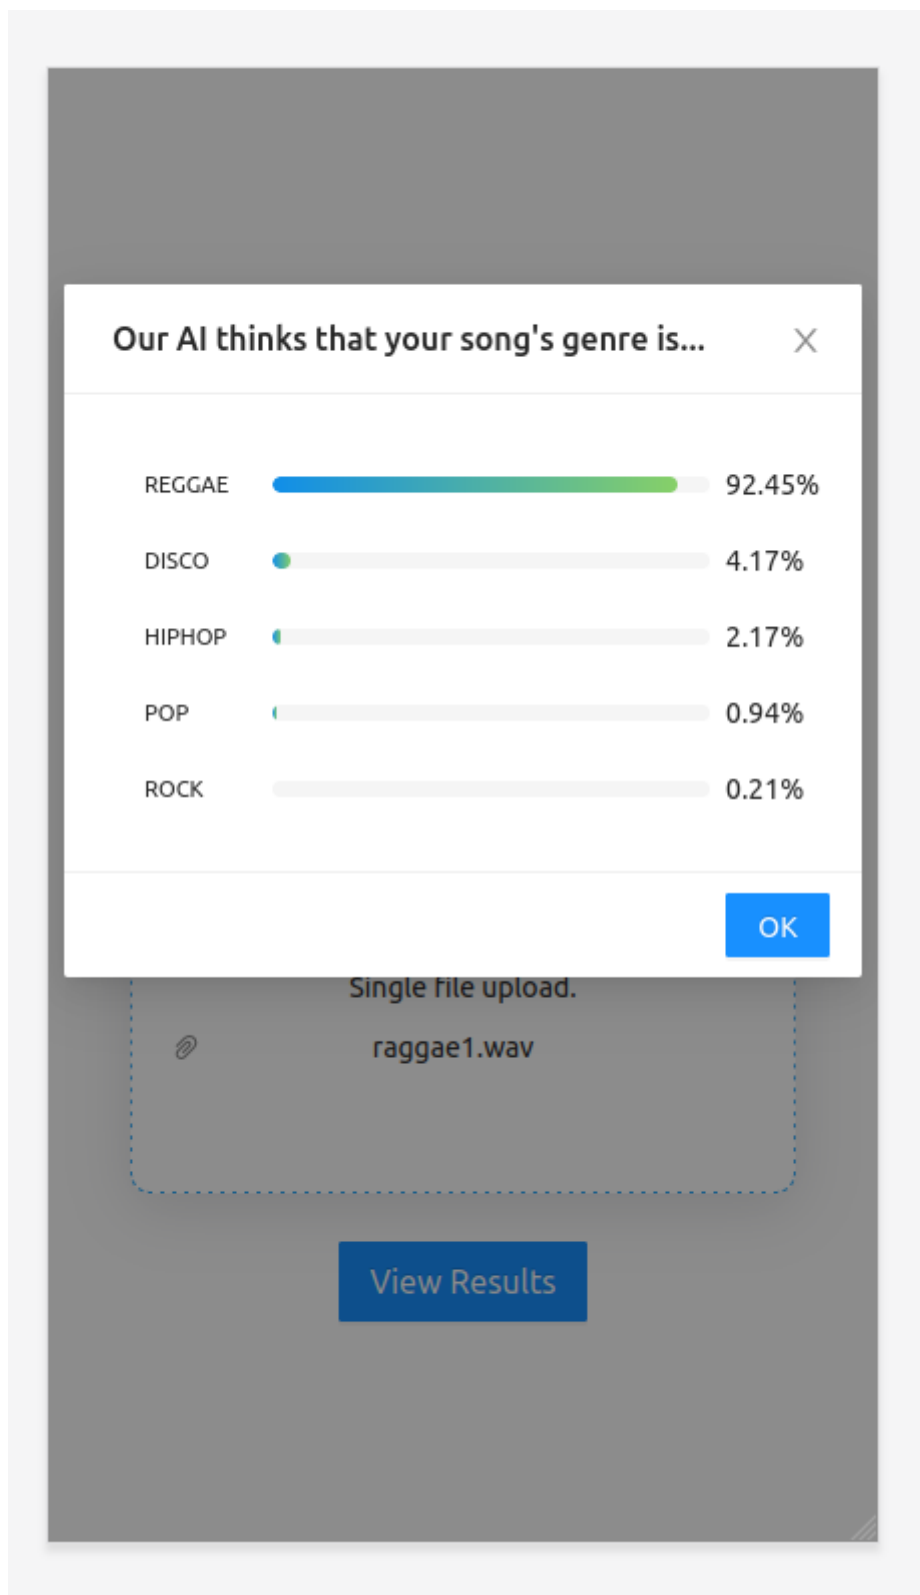


Figure 3: View prediction results

## Conclusion

For our project, we have used a model that uses both CNN and RNN in parallel. We take in .wav music audio files as input, and convert them into visual spectrograms before passing it to our model. Though our accuracy and f1 score we have obtained for our test set is not as high as we had hoped to be, when we compare our model with the state-of-the-art accuracy by the medium post (51%), we feel that it was indeed a good learning process, and the results we have obtained might be limited due to the aforementioned reasons.